

Lecture 16 - 12/11/2019

Container Security	2
Achieving Security	2
Secure Images	2
Secure Repositories	2
Secure deployment	3
Secure Runtime	3
Secure Orchestration	3
Secure the Host OS	3
Continuously monitor for security	3
Information Flow	4
Compiler based program protection	4

Container Security

Containers make achieving security much harder

- Shorter dev cycles
- More interactions
- Shared resources
- Public repositories are public
 - Anyone can put an image onto a public repo
- Container defaults can introduce vulnerabilities
 - E.g. "Expose 80"
- E.g. always specify a tag in the FROM statement
 - Don't just use :latest

Achieving Security

Secure Images

- Everything that you use in an image can be a vulnerability.
- Only pull in the stuff that you actually need
- Don't start with a base image from a public repo
 - Start from scratch if possible
- If you must start from a base image, only use a base that has been signed.
 - And from a trusted source
 - Scan for vulnerabilities before using.

Secure Repositories

- Continuously monitor
- If possible, use a private repository
 - Access control
- Useful security features can be added to containers in that repo
 - You can add image metadata which is useful for tracking vulnerabilities
 - Tagging to filter or sort images
 - Automated policy checking

Secure deployment

- Ensure that all containers build on top of each other
- Immutable container states

Secure Runtime

- Establish baseline behaviour for container in a normal, secure state.
- Network microservices: attach surface is large and complex
 - Allow only connectivity between container that actually need it
 - Restrict open ports and who can use them

Secure Orchestration

- Prevent risks from over-privileged accounts
- Prevent risks from attacks over the network
- Prevent risks from unwanted lateral movement.
- Configure orchestration to use proper access control
- Least privilege for each container
- White listing for specific containers

Secure the Host OS

- Scan for vulnerabilities
- Harden according to relevant guidelines/benchmarks
- Ensure container isolation

Continuously monitor for security

- Log every access to containers apps, services, systems, e.t.c.
- Performing regular audits of your log files
- Monitoring for anomalies
- Stay on top of current research

Information Flow

- Flow of information through a system
 - Confidentiality
 - Data of lower confidentiality flowing to a process of higher confidentiality
 - Integrity
 - Data of lower integrity flowing to a process of higher integrity
- This flow can happen two ways
 - Through code/programs
 - Compiler-based mechanisms to monitor and protect
 - Executable-based mechanisms (runtime mechanisms)
 - Through channels
 - System mechanisms to monitor and protect
 - Secure protocols to monitor and protect

Compiler based program protection

- Imagine code with two variables x and y
 - We can imagine a command sequence which might be several lines of code
 - Within this sequence it is defined that there is a flow of info from x to y if after this code we can look at y and infer something about x
- Explicit
 - $y := x;$
 - $tmp := y; y := tmp;$
- Implicit
 - $\text{If } (x = 1) \{ y = 0; \} \text{ else } \{ y = 1; \}$
- If x is a variable, then x is the "Information flow class" of that variable
 - Info can flow from x to y , if $\underline{x} \leq \underline{y}$ (confidentiality)
 - Info can flow from x to y , if $\underline{x} \geq \underline{y}$ (integrity)
- If there are several classes (e.g. $\underline{A}, \underline{B}, \underline{C}$)
 - I.e. least upper bound $\{\underline{A}, \underline{B}, \underline{C}\} \leq \underline{y}$
- Compiler based protection mechanisms checks that info flows through a program are authorization
- A set of program statements is certified with respect to an info flow policy if the info flows in these statements do not violate policy.
- E.g. Consider the statements
 - $\text{If } x = 1 \text{ then } y := a;$
 $\text{else } y := b$
 - Information flows to $\{x, a\}$ into y , or $\{x, b\}$ into y .
 - If $\underline{a} \leq \underline{y}$, $\underline{b} \leq \underline{y}$ and $\underline{x} \leq \underline{y}$ then the info flow is secure

- If the security depends on something such as time of day or some other unpredictable state, we can not verify with a compiler based mechanism.

Statements

- Assignment
- Compound
- Conditional
- Iterative
- Goto
- Proc
- Functions
- I/O statements

E.g. $y := f(x_1, x_2, x_3 \dots x_n)$

Least upper bound $\{\underline{x_1}, \underline{x_2}, \underline{x_3} \dots \underline{x_n}\} \leq y$

E.g.

```

If (x1 ... xn) then
    S1;
Else
    S2;

```

Check for the information flows within S1 and S2 and

Glb = greatest lower bound

Lub $\{\underline{x_1} \dots \underline{x_n}\} \leq \text{glb } \{\underline{y} \mid y \text{ is the target of some assignment statement in } s1 \text{ or } s2\}$

E.g. infinite loop

```

Proc SurpriseFlow (bool x, bool y) {
    y := 0;
    while x= 0 do
        (nothing)
    y := 1
}

```

X flows to y